

```
In [11]: # Imports
import pandas as pd
import numpy as np
import plotly as px
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

```
In [12]: df = pd.read_csv("/Users/jay/creditcardfraud/creditcard.csv")
```

```
In [65]: df.head()
```

```
Out[65]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.09
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.08
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.24
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.37
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.27

5 rows × 31 columns

```
In [66]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null float64
 1   V1          284807 non-null float64
 2   V2          284807 non-null float64
 3   V3          284807 non-null float64
 4   V4          284807 non-null float64
 5   V5          284807 non-null float64
 6   V6          284807 non-null float64
 7   V7          284807 non-null float64
 8   V8          284807 non-null float64
 9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

In [67]:

```
df.isnull().sum()
```

```
Out[67]: Time      0
         V1        0
         V2        0
         V3        0
         V4        0
         V5        0
         V6        0
         V7        0
         V8        0
         V9        0
         V10       0
         V11       0
         V12       0
         V13       0
         V14       0
         V15       0
         V16       0
         V17       0
         V18       0
         V19       0
         V20       0
         V21       0
         V22       0
         V23       0
         V24       0
         V25       0
         V26       0
         V27       0
         V28       0
         Amount    0
         Class     0
         dtype: int64
```

```
In [68]: df.describe()
```

```
Out[68]:
```

	Time	V1	V2	V3	V4	V5
<b>count</b>	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
<b>mean</b>	94813.859575	3.918649e-15	5.682686e-16	-8.761736e-15	2.811118e-15	-1.552100e-15
<b>std</b>	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380200e+00
<b>min</b>	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137400e+01
<b>25%</b>	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915000e-01
<b>50%</b>	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.430000e-02
<b>75%</b>	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119200e-01
<b>max</b>	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480100e+01

8 rows x 31 columns

```
In [20]: # legit and fraud transaction data
fraud = df[df['Class']==1]
legit = df[df['Class']==0]
```

```
In [21]: fraud['Amount'].describe()
```

```
Out[21]: count      492.000000
mean       122.211321
std        256.683288
min         0.000000
25%         1.000000
50%         9.250000
75%        105.890000
max        2125.870000
Name: Amount, dtype: float64
```

```
In [22]: legit['Amount'].describe()
```

```
Out[22]: count      284315.000000
mean           88.291022
std          250.105092
min           0.000000
25%           5.650000
50%          22.000000
75%          77.050000
max        25691.160000
Name: Amount, dtype: float64
```


Here the mean of legit transaction is \$88 and the mean of the fraudulent transaction is \$122 which is much higher than the legit transaction even if the max amount on legit transaction is \$25691

```
In [23]: sns.countplot('Class', data=df)
```

```
/Users/jay/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py
:36: FutureWarning: Pass the following variable as a keyword arg: x. From v
ersion 0.12, the only valid positional argument will be `data`, and passing
other arguments without an explicit keyword will result in an error or misi
nterpretation.
```

```
warnings.warn(
```

```
Out[23]: <AxesSubplot:xlabel='Class', ylabel='count'>
```



```
In [24]: df.groupby('Class').mean()
```

```
Out[24]:
```

	Time	V1	V2	V3	V4	V5	V6
Class							
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453	0.002419
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737

2 rows x 32 columns

As we can see there is a substantial difference between the mean of legit and fraud transaction in almost all the features.

```
In [50]: legit_sample = legit.sample(n=492) # which is equivalent to the fraudulen
```

```
In [51]: # Concatenate legit_sample and fraud


df_undersample = pd.concat([legit_sample, fraud], axis=0)
```

```
In [27]: sns.countplot('Class', data=df_undersample)
```

/Users/jay/opt/anaconda3/lib/python3.8/site-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[27]: <AxesSubplot:xlabel='Class', ylabel='count'>
```



Now as we have equal counts of the type of transaction, we can feed in the data for machine learning

```
In [28]: # Let's recheck the mean again if there is still any difference in mean va.

df_undersample.groupby('Class').mean()
```

Out[28]:

	Time	V1	V2	V3	V4	V5	V6	
Class								
0	97271.987805	0.077204	0.074780	-0.037623	0.106936	0.088720	-0.014933	0.026
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.56

2 rows × 32 columns

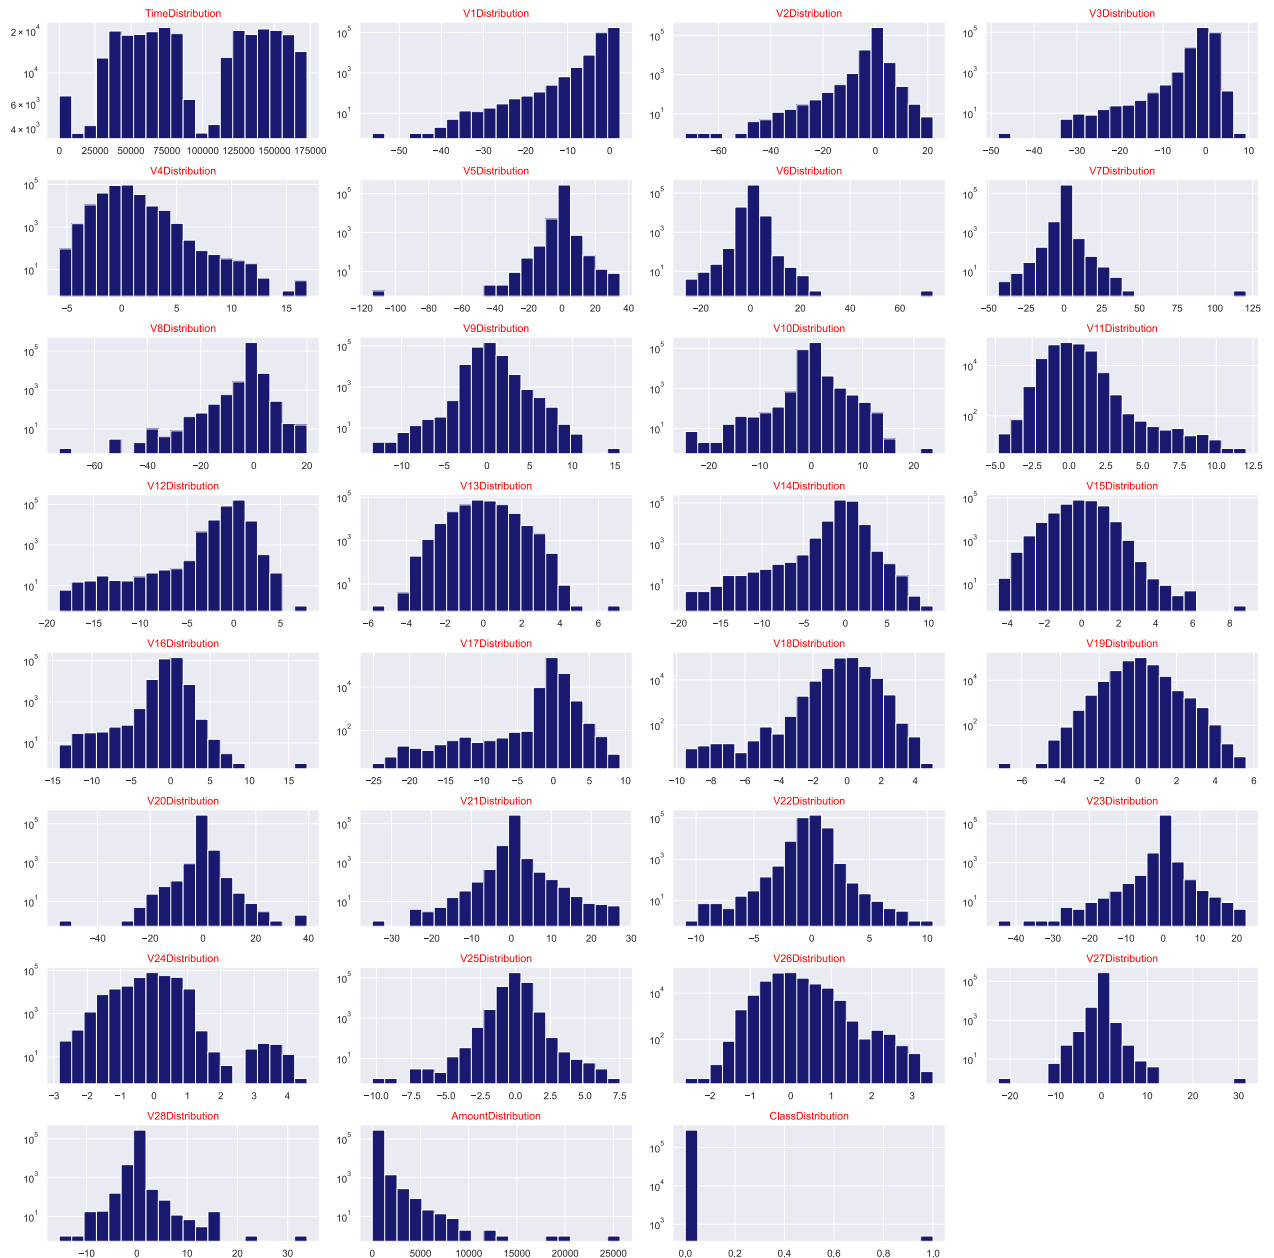
So even after discarding the major amount of data, there is still difference in mean between legit and fraud transaction

In [41]:

```
# Checking for the distribution

def multi_histogram(dataframe, features, rows, cols):
    fig = plt.figure(figsize=(20,20))
    for i, feature in enumerate(features):
        ax = fig.add_subplot(rows, cols, i+1)
        dataframe[feature].hist(bins=20, ax=ax, facecolor='midnightblue')
        ax.set_title(feature + "Distribution", color="red")
        ax.set_yscale('log')
    fig.tight_layout()
    plt.show()

multi_histogram(df, df.columns,8,4)
```



Since most of the data is already scaled, we need to scale the rest of the variables (Amount and time) and even the chart looks skewed to the right which means there are outliers or less bigger amount in the transaction history



```
In [29]: from sklearn.preprocessing import StandardScaler, RobustScaler

# Robust scaler can handle outliers
std_scaler = StandardScaler()
rob_scaler = RobustScaler()

df['scaled_amount'] = rob_scaler.fit_transform(df['Amount'].values.reshape(-1,))
df['scaled_time'] = rob_scaler.fit_transform(df['Time'].values.reshape(-1,))

df_scaled = df.drop(['Time', 'Amount'], axis=1,)
```

```
In [30]: df_scaled.head()
```

```
Out[30]:
```

	V1	V2	V3	V4	V5	V6	V7	V8
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533

5 rows × 31 columns

```
In [31]: df_scaled = df_scaled[['scaled_time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7',
                                'V8', 'V9', 'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17',
                                'V18', 'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28',
                                'scaled_amount', 'Class']
```

```
In [32]: df_scaled.head()
```

```
Out[32]:
```

	scaled_time	V1	V2	V3	V4	V5	V6	V7
0	-0.994983	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599
1	-0.994983	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803
2	-0.994972	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461
3	-0.994972	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
4	-0.994960	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941

5 rows × 31 columns

First we'll use under sampling method by taking the number of normal transaction equivalent to the fraudulent transaction.

```
In [70]: df_undersample = df_undersample.drop(['Amount', 'Time'], axis=1)
```

```
In [71]: df_undersample_scaled = df_undersample[['scaled_time', 'V1', 'V2', 'V3', 'V4',
        'V8', 'V9', 'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17',
        'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28',
        'scaled_amount', 'Class']
        ]]
```

```
In [72]: df_undersample_scaled.shape
```

```
Out[72]: (984, 31)
```

```
In [102... x = df_undersample_scaled.drop('Class', axis=1)
y = df_undersample_scaled['Class']
```

```
In [98]: x.shape, y.shape
```

```
Out[98]: ((984, 30), (984,))
```

```
In [99]: # Splitting data into training and testing data

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
```

```
In [103... print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(787, 30) (197, 30) (787,) (197,)
```

```
In [104... # Binary classification

from sklearn.linear_model import LogisticRegression

reg_log = LogisticRegression()

reg_log.fit(X_train, y_train)
```

Out[104... LogisticRegression()

In [105... `y_train_predict = reg_log.predict(X_train)`

In [109... *# Accuracy*

```
from sklearn.metrics import accuracy_score, classification_report

train_accuracy = accuracy_score( y_train, y_train_predict)
```

In [107... `print(train_accuracy)`

0.951715374841169

In [108... *# Now finding the accuracy on test set*

```
y_predict = reg_log.predict(X_test)

reg_accuracy_score = accuracy_score(y_test, y_predict)
print(reg_accuracy_score)
```

0.949238578680203

In [110... `print(classification_report(y_test, y_predict))`

	precision	recall	f1-score	support
0	0.91	0.99	0.95	95
1	0.99	0.91	0.95	102
accuracy			0.95	197
macro avg	0.95	0.95	0.95	197
weighted avg	0.95	0.95	0.95	197